

nimCSO: A Nim package for Compositional Space Optimization

Adam M. Krajewski¹, Arindam Debnath¹, Wesley F. Reinhart^{1,2},
Allison M. Beese¹, and Zi-Kui Liu¹

¹ Department of Materials Science and Engineering, The Pennsylvania State University, United States of America ² Institute for Computational and Data Sciences, The Pennsylvania State University, United States of America ¶ Corresponding author

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

Software

- [Review](#)
- [Repository](#)
- [Archive](#)

Editor: [Open Journals](#)

Reviewers:

- [@openjournals](#)

Submitted: 01 January 1970

Published: unpublished

License

Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License ([CC BY 4.0](#)).

Summary

nimCSO is a high-performance tool implementing several methods for selecting components (data dimensions) in compositional datasets, which optimize the data availability and density for applications such as machine learning. Making said choice is a combinatorically hard problem for complex compositions existing in high-dimensional spaces due to the interdependency of components being present. Such spaces are encountered across many scientific disciplines (see [Statement of Need](#)), including materials science, where datasets on Compositionally Complex Materials (CCMs) often span 20-45 chemical elements, 5-10 processing types, and several temperature regimes, for up to 60 total data dimensions. This challenge also exists in everyday contexts, such as study of cooking ingredients ([Ahn et al., 2011](#)), which interact in various recipes, giving rise to questions like “*Given 100 spices at the supermarket, which 20, 30, or 40 should I stock in my pantry to maximize the number of unique dishes I can spice according to recipe?*”. Critically, this is not as simple as frequency-based selection because, e.g., removing less common nutmeg and cinnamon from your shopping list will prevent many recipes with the frequent vanilla, but won’t affect those using black pepper ([Edmisten, 2022](#)).

At its core, nimCSO leverages the metaprogramming ability of the Nim language ([Rumpf, 2023](#)) to optimize itself at compile time, both in terms of speed and memory handling, to the specific problem statement and dataset at hand based on a human-readable configuration file. As demonstrated in the [Methods and Performance](#) section, nimCSO reaches the physical limits of the hardware (L1 cache latency) and can outperform an efficient native Python implementation over 100 times in terms of speed and 50 times in terms of memory usage (*not* counting the interpreter), while also outperforming the NumPy ([Harris et al., 2020](#)) based implementation 37 and 17 times, respectively, when checking a candidate solution.

nimCSO is designed to be both (1) a user-ready tool, implementing two efficient brute-force approaches (for handling up to 25 dimensions), a custom search algorithm (for up to 40 dimensions), and a genetic algorithm (for any dimensionality), and (2) a scaffold for building even more elaborate methods in the future, including heuristics going beyond data availability. All configuration is done with a simple human-readable YAML file and plain text data files, making it easy to modify the search method and its parameters with no knowledge of programming and only basic command line skills.

Statement of Need

nimCSO is an interdisciplinary tool applicable to any field where data is composed of a large number of independent components and their interaction is of interest in a modeling effort, ranging from economics where factor selection affects performance of analytical ([Fan et al.,](#)

2013) and ML (Peng et al., 2021) models, through medicine where drug interactions can have a significant impact on the treatment (Maher et al., 2014) (an escalating problem (Guthrie et al., 2015)) and understanding of microbial interactions can help fight gastrointestinal problems (Berg et al., 2022; Leeuwen et al., 2023), to materials science, where the composition and processing history are critical to resulting properties. The latter has been the root motivation for the development of nimCSO within the ULTERA Project (ultera.org) carried under the US DOE ARPA-E ULTIMATE program, which aims to develop a new generation of ultra-high temperature materials for aerospace applications, through generative machine learning models (Debnath et al., 2021) driving thermodynamic modeling, alloy design, and manufacturing (Li et al., 2024).

One of the most promising materials for such applications are the aforementioned CCMs and their metal-focused subset of Refractory High Entropy Alloys (RHEAs) (Senkov et al., 2018), which have rapidly grown since first proposed by Cantor et al. (2004) and Yeh et al. (2004). Contrary to most of the traditional alloys, they contain many chemical elements (typically 4-9) in similar proportions in the hope of thermodynamically stabilizing the material by increasing its configurational entropy ($\Delta S_{conf} = \sum_i^N x_i \ln x_i$ for ideal mixing of N elements with fractions x_i), which encourages sampling from a large palette of chemical elements. At the time of writing, the ULTERA Database is the largest collection of HEA data, containing over 7,000 points manually extracted from 560 publications. It covers 37 chemical elements resulting in extremely large compositional spaces (Krajewski et al., 2024); thus, it becomes critical to answer questions like “Which combination of how many elements will unlock the most expansive and simultaneously dense dataset?” which has $2^{37} - 1$ or 137 billion possible solutions.

Another significant example of intended use is to perform similar optimizations over large (many millions) datasets of quantum mechanics calculations spanning 93 chemical elements and accessible through the OPTIMADE API (Evans et al., 2024).

Methods and Performance

Overview

As shown in Figure 1, nimCSO can be used as a user-tool based on human-readable configuration and a data file containing data “elements” which can be any strings representing problem-specific names of, e.g., spices, market stocks, drug names, or chemical systems. A single command is used to recompile (nim c -f) and run (-r) problem (-d:configPath=config.yaml) with nimCSO (src/nimcso) using one of several methods. Advanced users can also quickly customize the provided methods with brief scripts using the nimCSO as a data-centric library.

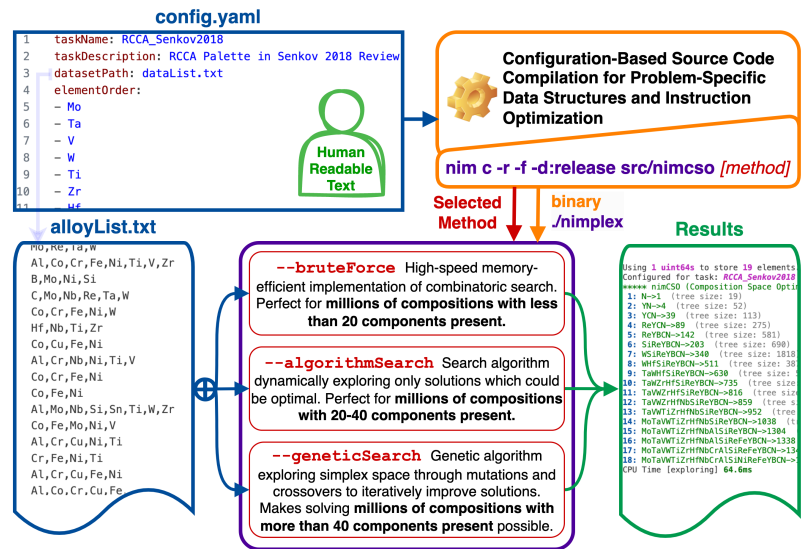


Figure 1: Schematic of core nimCSO data flow with a description of key methods. Metaprogramming is used to compile the software optimized to the human-readable data and configuration files at hand.

Internally, nimCSO is built around storing the data and solutions in one of two ways. The first is as bits inside an integer (uint64), which allows for the highest speed and lowest memory consumption possible but is limited to 64 dimensions and does not allow for easy extension to other use cases; thus, as of publication, it is used only in a particular bruteForceInt routine. The second one, used in bruteForce, algorithmSearch, and geneticSearch, implements a custom easily extensible ElSolution type containing heuristic value and BitArray payload, which is defined at compile time based on the configuration file to minimize necessary overheads. As shown in Table 1, both encodings outperform typical native Python and NumPy implementations, typically used in this application by the scientific community, which can be found in benchmarks directory.

Table 1: Benchmarks of (1) the average time to evaluate how many datapoints would be lost if 5 selected components were removed from a dataset with 2,150 data points spanning 37 components, averaged over 10,000 runs, and (2) the size of the data structure representing the dataset. Values were obtained by running scripts in benchmarks directory on Apple M2 Max CPU. Pre-processing time is excluded, as it has negligible impact on larger, realistic problems.

Tool	Object	Time per Dataset	Time per Entry (Relative)	Database Size (Relative)
Python ^{3.11}	set	107.5 μ s	50.0 ns ($\times 1$)	871.5 kB ($\times 1$)
NumPy ^{1.26}	array	36.4 μ s	16.9 ns ($\times 3.0$)	79.7 kB ($\times 10.9$)
nimCSO ^{0.6}	BitArray	6.9 μ s	3.2 ns ($\times 15.6$)	50.4 kB ($\times 17.3$)
nimCSO ^{0.6}	uint64	0.98 μ s	0.456 ns ($\times 110$)	16.8 kB ($\times 52$)

Brute-Force Search

The brute-force search is a naïve method of evaluating all possibilities; however, its near-zero overhead can make it the most efficient for small problems. In this implementation, all entries in the *power set* of N considered elements are represented as a range of integers from 0 to $2^N - 1$, and used to initialize uint64/BitArrays on the fly. To minimize the memory footprint of solutions, the algorithm only keeps track of the best solution for a given number of elements present in the solution. Current implementations are limited to 64 elements, as it is not

feasible beyond approximately 30 elements; however, the one based on BitArray could be easily extended if needed.

Algorithm-Based Search

The algorithm implemented in the `algorithmSearch` routine, targeting high dimensional problems (20-50), iteratively expands and evaluates candidates from a priority queue (implemented through an efficient binary heap (Williams, 1964)) while leveraging the fact that *the number of data points lost when removing elements A and B from the dataset has to be at least as large as when removing either A or B alone* to delay exploration of candidates until they can contribute to the solution. Furthermore, to (1) avoid revisiting the same candidate without keeping track of visited states and to (2) avoid exhaustive search by inhibiting the exploration of unlikely candidates, the algorithm *assumes* that while searching for a given order of solution, elements present in already expanded solutions will not improve those not yet expanded. This effectively prunes candidate branches requiring two or more levels of backtracking, making the algorithm computationally feasible in higher-dimensional problems at the cost of potentially suboptimal results. In the authors' tests over HEAs, this algorithm has generated correct results (agreeing with an exhaustive bruteForce exploration), except for occasional differences in the last explored solution, where almost all data was discarded, and a few remaining points were highly independent.

Genetic Search

Beyond 50 components, the `algorithm-based` method will likely run out of memory on most personal systems. The `geneticSearch` routine resolves this issue through an evolution strategy to iteratively improve solutions based on custom mutate and crossover procedures. Both are of uniform type (Goldberg, 1989) with additional constraint of Hamming weight (Knuth, 2009) preservation in order to preserve number of considered elements in parents and offspring. In mutate this is achieved by using purely random bit swapping, rather than more common flipping, as demonstrated in the Figure 2.




Figure 2: Schematic of mutate procedure where bits are swapping randomly, so that (1) bit can swap itself, (2) bits can swap causing a flip, or (3) bits can swap with no effect.

Meanwhile, in crossover, this constraint is satisfied by passing overlapping bits directly, while non-overlapping bits are shuffled and distributed at positions present in one of the parents, as shown in Figure 3.




Figure 3: Schematic of uniform crossover procedure preserving Hamming weight implemented in `nimCSO`.

The above are applied iteratively, with best solutions carried to next generation, until the solution converges or the maximum number of iterations is reached. Unlike the other methods,

124 the present method is not limited by the number of components and lets user control both
125 time and memory requirements, either to make big problems feasible or to get a good-enough
126 solution quickly in small problems. However, it comes with no optimality guarantees.

127 **Use Examples**

128 The tool comes with two pre-defined example problems to demonstrate its use. The first one is
129 defined in the default config.yaml file and goes through the complete dataset of 2,150 data
130 points spanning 37 components in dataList.txt based on the ULTERA Dataset (Debnath et
131 al., 2021). It is intended to showcase algorithmSearch/-as and geneticSearch/-gs methods,
132 as brute-forcing would take around one day. The second one is defined in config_rhea.yaml
133 and uses the same dataset but a limited scope of components critical to RHEAs (Senkov et
134 al., 2018) and is intended to showcase bruteForce/-bf and bruteForceInt/-bfi methods.
135 With four simple commands (see Table 2), the user can compare the methods' performance
136 and the solutions' quality.

Table 2: Four example tasks alongside typical CPU time and memory usage on Apple M2 Max.

Task Definition (nim c -r -f -d:release ...)	Time (s)	Memory (MB)
-d:configPath=config.yaml src/nimcso -as	308s	488 MB
-d:configPath=config.yaml src/nimcso -gs	5.10s	3.2 MB
-d:configPath=config_rhea.yaml src/nimcso -as	0.073s	2.2 MB
-d:configPath=config_rhea.yaml src/nimcso -gs	0.426s	2.1 MB
-d:configPath=config_rhea.yaml src/nimcso -bf	3.726s	2.0 MB
-d:configPath=config_rhea.yaml src/nimcso -bfi	0.495s	2.0 MB

137 In case of issues, the help message can be accessed by running the tool with -h flag or by
138 referring to documentation at amkrajewski.github.io/nimCSO.

139 **Contributions**

140 A.M.K. was responsible for conceptualization, methodology, software, testing and validation,
141 writing of manuscript, and visualization; A.D. was responsible for testing software and results
142 in training machine learning models; A.M.B., W.F.R., Z-K.L. were responsible for funding
143 acquisition, review, and editing. Z-K.L. was also supervising the work.

144 **Acknowledgements**

145 This work has been funded through grants: DOE-ARPA-E DE-AR0001435, NSF-POSE
146 FAIN-2229690, and ONR N00014-23-2721. We would also like to acknowledge Dr. Jonathan
147 Siegel at Texas A&M University for several valuable discussions and feedback on the project,
148 and Dr. Bradley Dice for his code contributions to the project during the review process, such
149 as optimization of the native Python implementation.

150 **References**

151 Ahn, Y.-Y., Ahnert, S. E., Bagrow, J. P., & Barabási, A.-L. (2011). Flavor network and the
152 principles of food pairing. *Scientific Reports*, 1, 196. <https://doi.org/10.1038/srep00196>
153 Berg, N. I. van den, Machado, D., Santos, S., Rocha, I., Chacón, J., Harcombe, W., Mitri,
154 S., & Patil, K. R. (2022). Ecological modelling approaches for predicting emergent

- properties in microbial communities. *Nature Ecology & Evolution*, 6, 855–865. <https://doi.org/10.1038/s41559-022-01746-7>
- Cantor, B., Chang, I. T. H., Knight, P., & Vincent, A. J. B. (2004). Microstructural development in equiatomic multicomponent alloys. *Materials Science and Engineering A*, 375–377, 213–218. <https://doi.org/10.1016/j.msea.2003.10.257>
- Debnath, A., Krajewski, A. M., Sun, H., Lin, S., Ahn, M., Li, W., Priya, S., Singh, J., Shang, S., Beese, A. M., Liu, Z.-K., & Reinhart, W. F. (2021). *Journal of Materials Informatics*, 1. <https://doi.org/10.20517/jmi.2021.05>
- Edmisten, W. (2022). *Analyzing all recipes*. <https://wcedmisten.fyi/post/analyzing-all-recipes/#spices>
- Evans, M. L., Bergsma, J., Merkys, A., Andersen, C. W., Andersson, O. B., Beltrán, D., Blokhin, E., Boland, T. M., Balderas, R. C., Choudhary, K., Díaz, A. D., García, R. D., Eckert, H., Eimre, K., Montero, M. E. F., Krajewski, A. M., Mortensen, J. J., Duarte, J. M. N., Pietryga, J., ... Armiento, R. (2024). *Developments and applications of the OPTIMADE API for materials discovery, design, and data exchange*. <https://doi.org/10.48550/arXiv.2402.00572>
- Fan, J., Liao, Y., & Mincheva, M. (2013). Large covariance estimation by thresholding principal orthogonal complements. *Journal of the Royal Statistical Society Series B: Statistical Methodology*, 75, 603–680. <https://doi.org/10.1111/rssb.12016>
- Goldberg, D. E. (1989). *Genetic algorithms in search, optimization and machine learning* (1st ed.). Addison-Wesley Longman Publishing Co., Inc. ISBN: 0201157675
- Guthrie, B., Makubate, B., Hernandez-Santiago, V., & Dreischulte, T. (2015). The rising tide of polypharmacy and drug-drug interactions: Population database analysis 1995–2010. *BMC Medicine*, 13, 74. <https://doi.org/10.1186/s12916-015-0322-7>
- Harris, C. R., Millman, K. J., Walt, S. J. van der, Gommers, R., Virtanen, P., Cournapeau, D., Wieser, E., Taylor, J., Berg, S., Smith, N. J., Kern, R., Picus, M., Hoyer, S., Kerkwijk, M. H. van, Brett, M., Haldane, A., Río, J. F. del, Wiebe, M., Peterson, P., ... Oliphant, T. E. (2020). Array programming with NumPy. *Nature*, 585(7825), 357–362. <https://doi.org/10.1038/s41586-020-2649-2>
- Knuth, D. E. (2009). *The art of computer programming, volume 4, fascicle 1: Bitwise tricks & techniques; binary decision diagrams* (12th ed.). Addison-Wesley Professional. ISBN: 0321580508
- Krajewski, A. M., Beese, A. M., Reinhart, W. F., & Liu, Z.-K. (2024). *Efficient generation of grids and traversal graphs in compositional spaces towards exploration and path planning exemplified in materials*. <https://doi.org/10.48550/arXiv.2402.03528>
- Leeuwen, P. T. van, Brul, S., Zhang, J., & Wortel, M. T. (2023). Synthetic microbial communities (SynComs) of the human gut: Design, assembly, and applications. *FEMS Microbiology Reviews*, 47. <https://doi.org/10.1093/femsre/fuad012>
- Li, W., Raman, L., Debnath, A., Ahn, M., Lin, S., Krajewski, A. M., Shang, S., Priya, S., Reinhart, W. F., Liu, Z.-K., & Beese, A. M. (2024). Design and validation of refractory alloys using machine learning, CALPHAD, and experiments. *International Journal of Refractory Metals and Hard Materials*, 121, 106673. <https://doi.org/10.1016/j.ijrmhm.2024.106673>
- Maher, R. L., Hanlon, J., & Hajjar, E. R. (2014). Clinical consequences of polypharmacy in elderly. *Expert Opinion on Drug Safety*, 13, 57–65. <https://doi.org/10.1517/14740338.2013.827660>
- Peng, Y., Albuquerque, P. H. M., Kimura, H., & Saavedra, C. A. P. B. (2021). Feature selection and deep neural networks for stock price direction forecasting using technical analysis indicators. *Machine Learning with Applications*, 5, 100060. <https://doi.org/10.1016/j.mlwa.2021.100060>

- 203 Rumpf, A. (2023). *Nim programming language v2.0.0*. <https://nim-lang.org/>
- 204 Senkov, O. N., Miracle, D. B., Chaput, K. J., & Couzinie, J.-P. (2018). Development and
205 exploration of refractory high entropy alloys—a review. *Journal of Materials Research*, 33,
206 3092–3128. <https://doi.org/10.1557/jmr.2018.153>
- 207 Williams, J. W. J. (1964). Algorithm 232 - heapsort. *Communications of the ACM*, 7, 347–349.
208 <https://doi.org/10.1145/512274.512284>
- 209 Yeh, J. W., Chen, S. K., Lin, S. J., Gan, J. Y., Chin, T. S., Shun, T. T., Tsau, C. H., &
210 Chang, S. Y. (2004). Nanostructured high-entropy alloys with multiple principal elements:
211 Novel alloy design concepts and outcomes. *Advanced Engineering Materials*, 6, 299–303.
212 <https://doi.org/10.1002/adem.200300567>

DRAFT